


(a) security classification, PROMISE_exp — the same 968 items scored in all three regimes

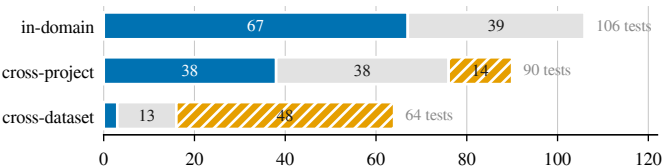
■ BERT_w □ RoBERTa_w ▨ 8 prompted LLMs (median)



increasing distribution shift between what the encoder was fitted on and what it is scored on →

(b) every one of the 260 paired comparisons, grouped by regime

■ Encoder wins ■ n.s. ▨ LLM wins



number of paired encoder-vs-LLM tests falling each way (exact McNemar, Benjamini-Hochberg, $\alpha = 0.05$)